

Computer Vision: Image filtering

Siheng Chen 陈思衡

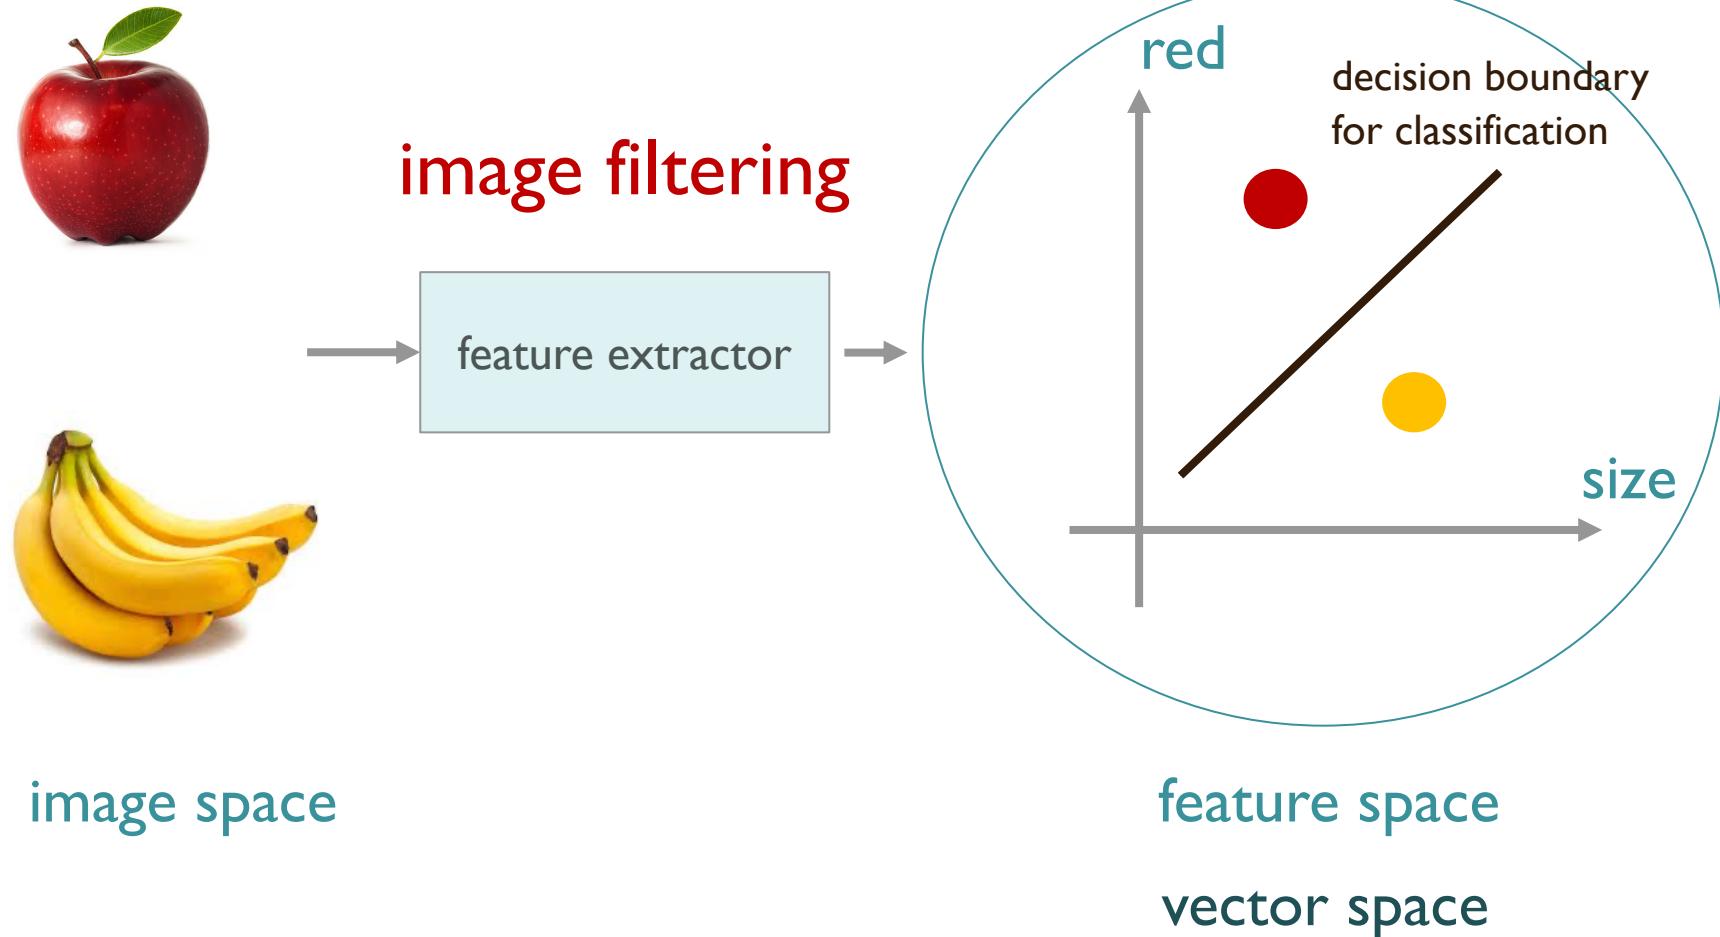
Image filtering

Intro to image filtering

More maths about filtering

Image filter design

Why we talk about image filtering



Why we talk about image filtering?

Image editing



Why we talk about image filtering?

Image denoising



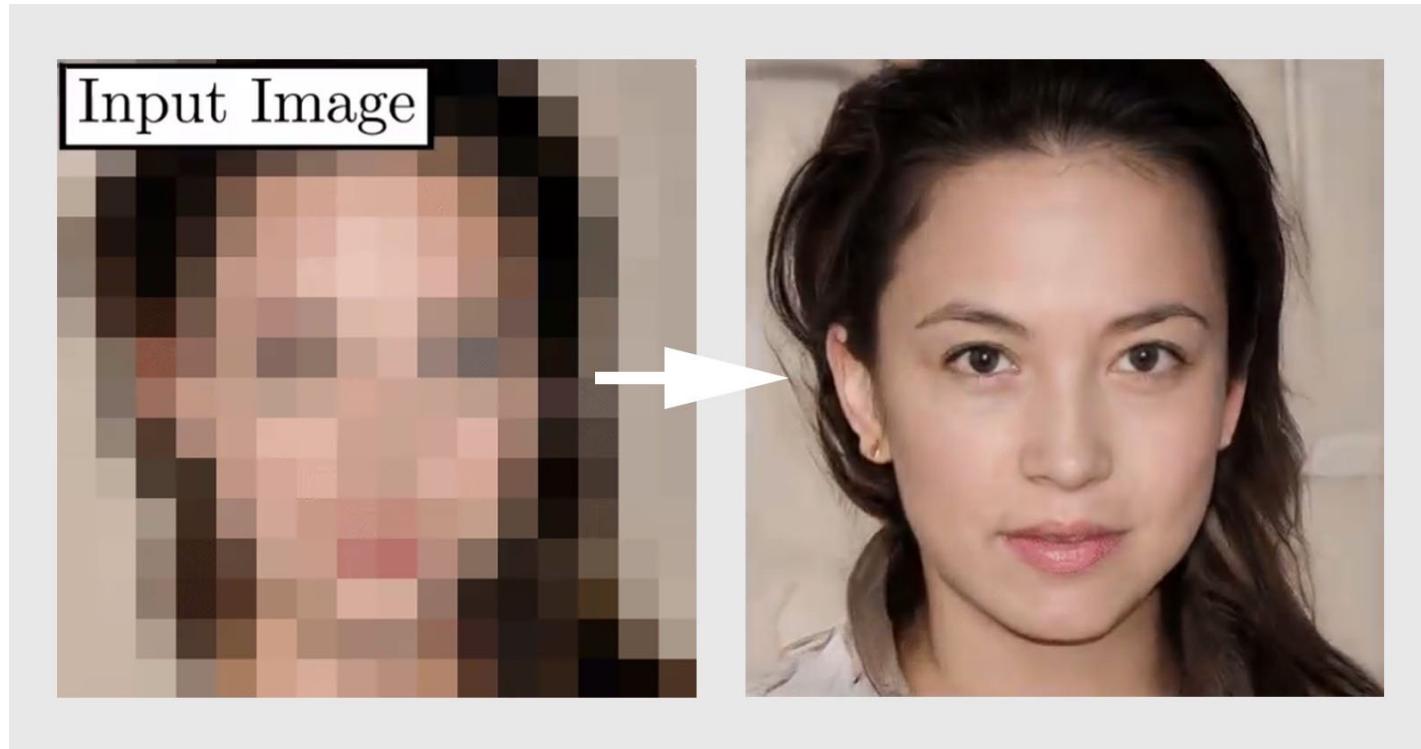
Why we talk about image filtering?

Image inpainting



Why we talk about image filtering?

Image super-resolution



What is an image?

A (color) image is a 3D tensor of numbers.



What is an image?

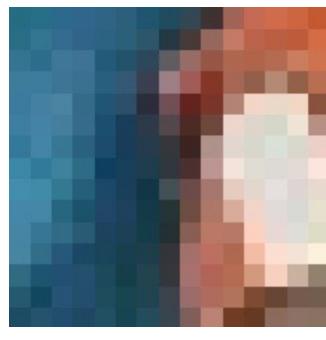
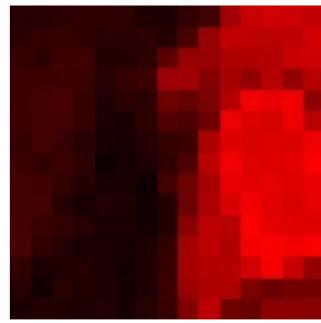
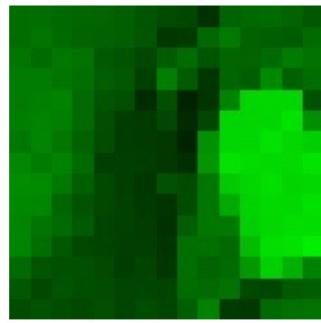


image patch

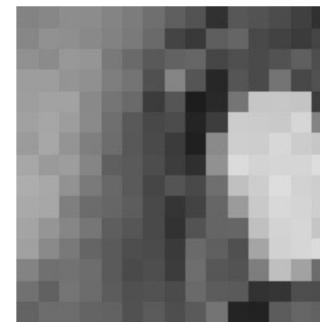
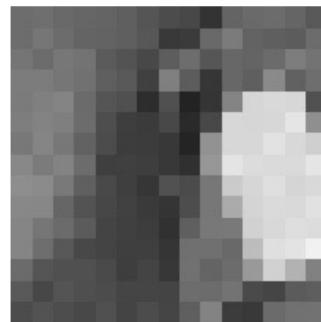
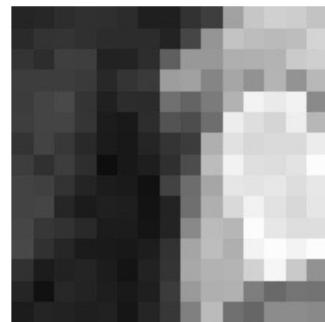
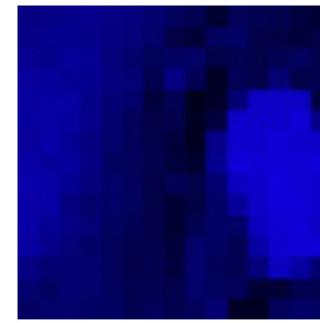
red



green



blue



Intensity value at each channel

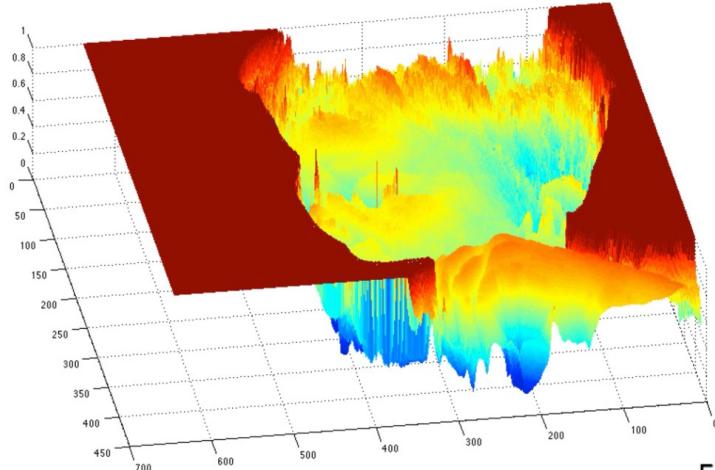
Each channel is a 2D array of numbers

What is an image?

A (grayscale) image is a 2D function.



$$f(\mathbf{x})$$

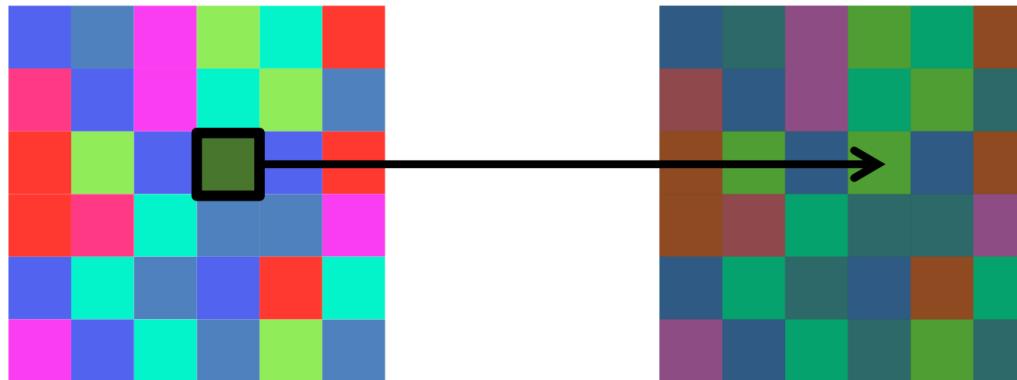


$$\text{pixel } \mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$$

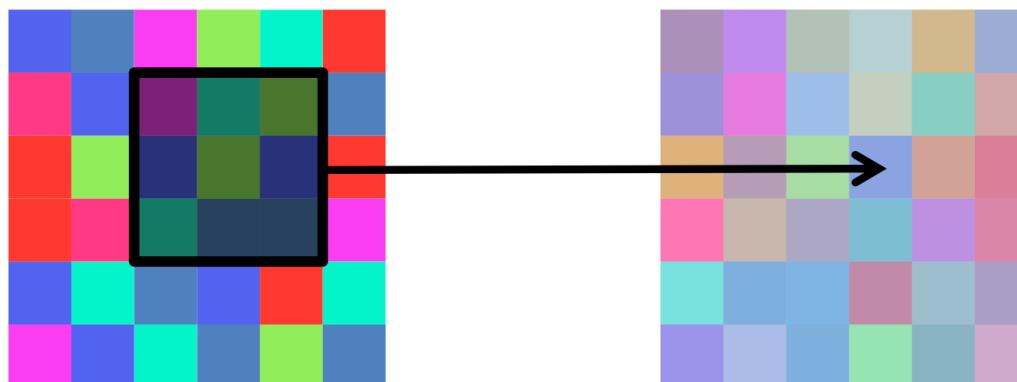
Image filtering is an operation to modify the 2D function

Image filtering

Point-wise operation



Neighbor-wise operation



Point-wise operation

original



darken



lower contrast

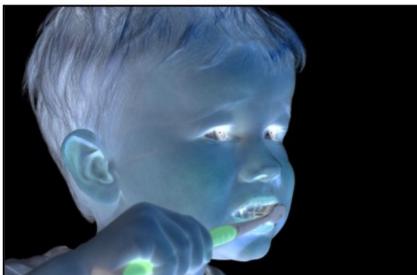


non-linear lower contrast

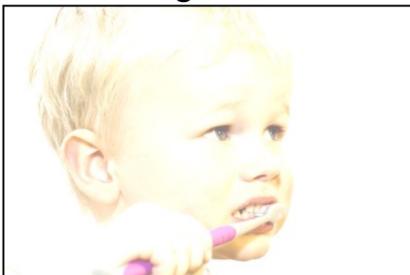


x

invert



lighten



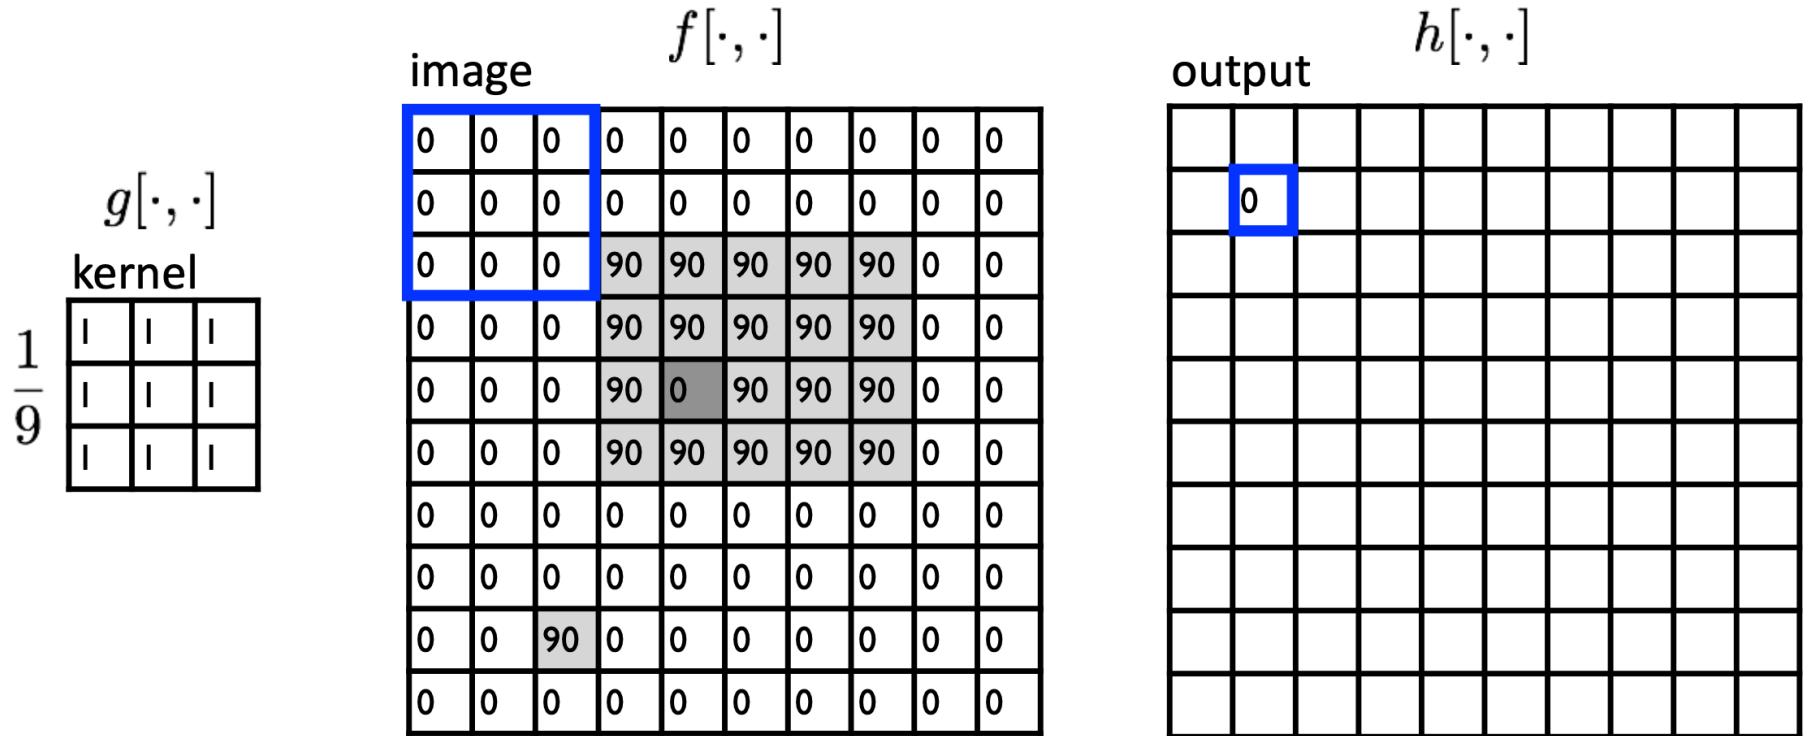
raise contrast



non-linear raise contrast



Neighbor-wise operation



Neighbor-wise operation

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

image

$f[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

output

$h[\cdot, \cdot]$

			0	10						

Neighbor-wise operation

input



filter

0	0	0
0	1	0
0	0	0

output



unchanged

Neighbor-wise operation

input



filter

0	0	0
0	0	1
0	0	0

output



shift to left by one

Neighbor-wise operation

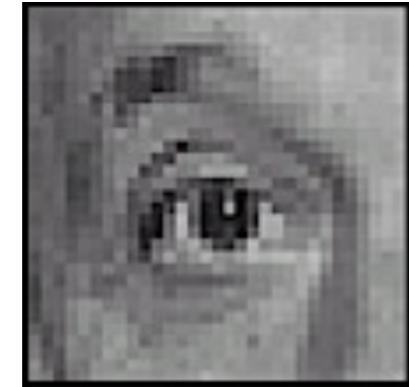
input



filter

$$\begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 0 & 2 & 0 \\ \hline 0 & 0 & 0 \\ \hline \end{array} - \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

output



sharpening

Image filtering

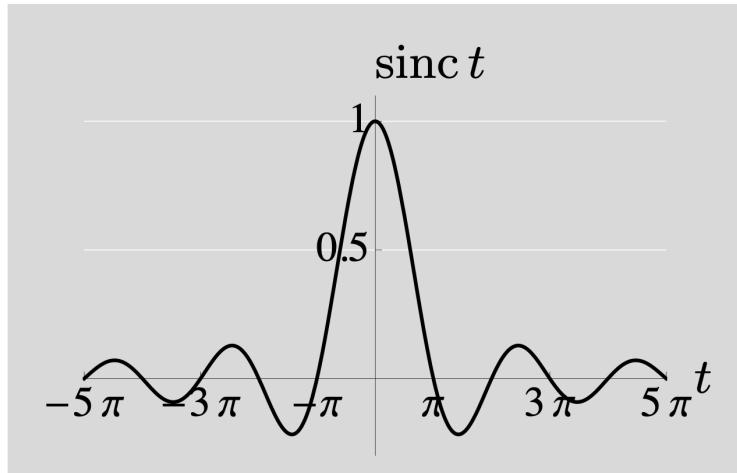
Intro to filtering

More maths about filtering

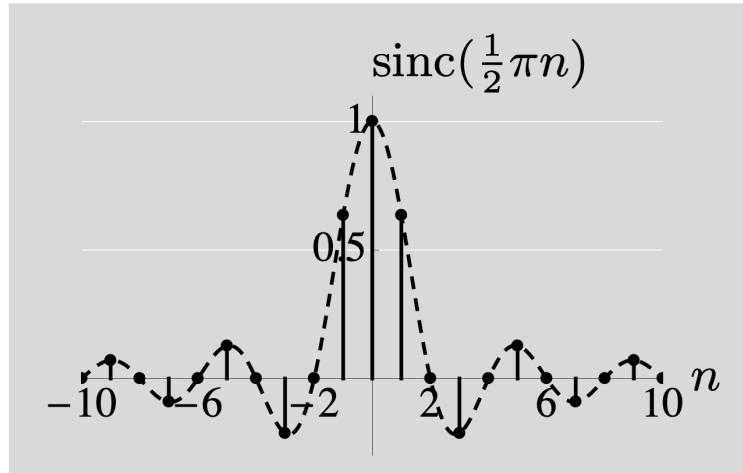
Image filter design

Filtering

Sequence



(a)



(b)

(a) The sinc function $\text{sinc } t$. (b) The sinc sequence $\text{sinc}(\frac{1}{2}\pi n)$.

Filtering

Sequence

- two-sided infinite

$$x = [\dots \ x_{-2} \ x_{-1} \boxed{x_0} \ x_1 \ x_2 \ \dots]^T$$



- a sequence of length N

$$x = [\boxed{x_0} \ x_1 \ x_2 \ \dots \ x_{N-1}]^T$$

Filtering

Sequence

- Autocorrelation

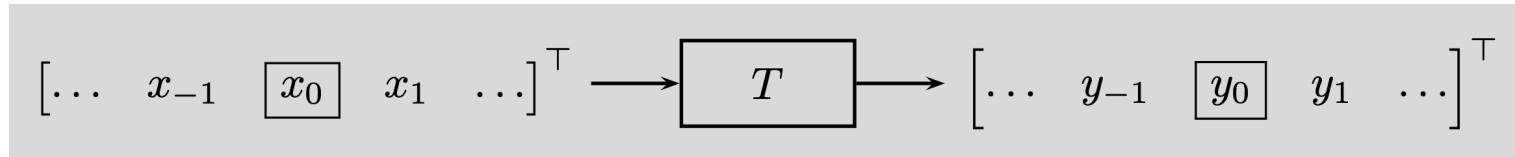
$$a_n = \sum_{k \in \mathbb{Z}} x_k x_{k-n}^* = \langle x_k, x_{k-n} \rangle_k$$

- Cross-correlation

$$c_n = \sum_{k \in \mathbb{Z}} x_k y_{k-n}^* = \langle x_k, y_{k-n} \rangle_k$$

Filtering

Systems



A discrete-time system

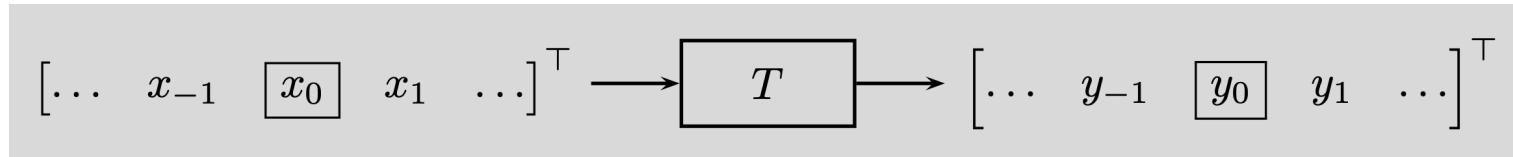
- Example

- The shift-by-1 operator, or delay $y_n = x_{n-1}$, $n \in \mathbb{Z}$

$$y = \begin{bmatrix} \vdots \\ y_{-1} \\ \boxed{y_0} \\ y_1 \\ \vdots \end{bmatrix} = \begin{bmatrix} \vdots \\ x_{-2} \\ \boxed{x_{-1}} \\ x_0 \\ \vdots \end{bmatrix} = \begin{bmatrix} \vdots & \vdots & \vdots & \vdots \\ \cdots & 0 & 0 & 0 & \cdots \\ \cdots & 1 & \boxed{0} & 0 & \cdots \\ \cdots & 0 & 1 & 0 & \cdots \\ \vdots & \vdots & \vdots & \vdots \end{bmatrix} \begin{bmatrix} \vdots \\ x_{-1} \\ \boxed{x_0} \\ x_1 \\ \vdots \end{bmatrix}$$

Filtering

Systems



A discrete-time system

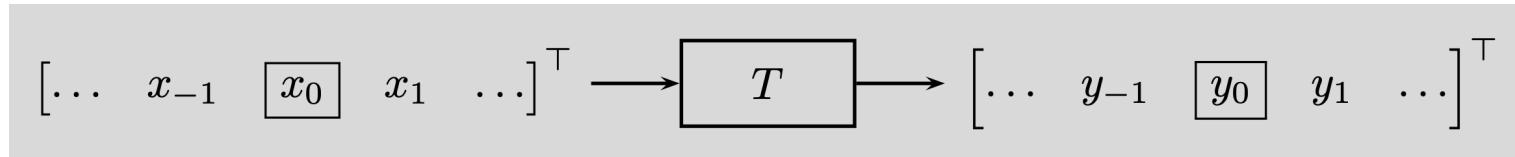
- Example

- The advance-by-1 operator $y_n = x_{n+1}, n \in \mathbb{Z}$

$$y = \begin{bmatrix} \vdots \\ y_{-1} \\ \boxed{y_0} \\ y_1 \\ \vdots \end{bmatrix} = \begin{bmatrix} \vdots \\ x_0 \\ \boxed{x_1} \\ x_2 \\ \vdots \end{bmatrix} = \begin{bmatrix} \vdots & \vdots & \vdots & \vdots \\ \dots & 0 & 1 & 0 & \dots \\ \dots & 0 & \boxed{0} & 1 & \dots \\ \dots & 0 & 0 & 0 & \dots \\ \vdots & \vdots & \vdots & \vdots \end{bmatrix} \begin{bmatrix} \vdots \\ x_{-1} \\ \boxed{x_0} \\ x_1 \\ \vdots \end{bmatrix}$$

Filtering

Systems



A discrete-time system

- Example

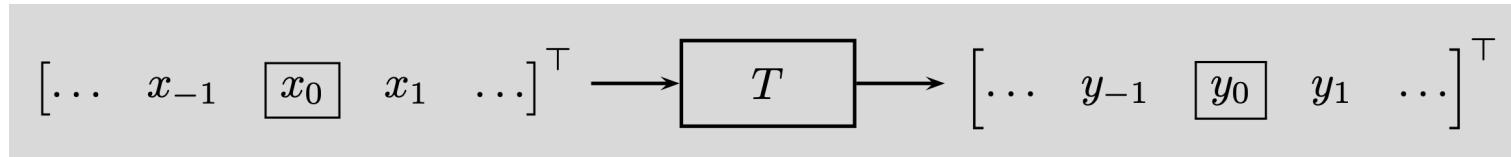
- Modulator

$$y_n = \alpha_n x_n, \quad n \in \mathbb{Z}$$

$$\begin{bmatrix} \vdots \\ y_{-1} \\ \boxed{y_0} \\ y_1 \\ \vdots \end{bmatrix} = \begin{bmatrix} \vdots \\ \alpha_{-1}x_{-1} \\ \boxed{\alpha_0x_0} \\ \alpha_1x_1 \\ \vdots \end{bmatrix} = \begin{bmatrix} \cdots & \alpha_{-1} & 0 & 0 & \cdots \\ \cdots & 0 & \boxed{\alpha_0} & 0 & \cdots \\ \cdots & 0 & 0 & \alpha_1 & \cdots \\ \vdots & \vdots & \vdots & \vdots & \vdots \end{bmatrix} \begin{bmatrix} \vdots \\ x_{-1} \\ \boxed{x_0} \\ x_1 \\ \vdots \end{bmatrix}$$

Filtering

Systems



A discrete-time system

- Example

- Averaging operators

$$y_n = \frac{1}{3}(x_{n-1} + x_n + x_{n+1}), \quad n \in \mathbb{Z}$$

$$\begin{bmatrix} \vdots \\ y_{-1} \\ \boxed{y_0} \\ y_1 \\ \vdots \end{bmatrix} = \frac{1}{3} \begin{bmatrix} \dots & \vdots & \vdots & \vdots & \vdots & \vdots \\ \dots & 1 & 1 & 1 & 0 & 0 & \dots \\ \dots & 0 & 1 & \boxed{1} & 1 & 0 & \dots \\ \dots & 0 & 0 & 1 & 1 & 1 & \dots \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \end{bmatrix} \begin{bmatrix} \vdots \\ x_{-2} \\ x_{-1} \\ \boxed{x_0} \\ x_1 \\ x_2 \\ \vdots \end{bmatrix}$$

Filtering

Systems

DEFINITION 3.1 (LINEAR SYSTEM) A discrete-time system T is called *linear* when, for any inputs x and y and any $\alpha, \beta \in \mathbb{C}$,

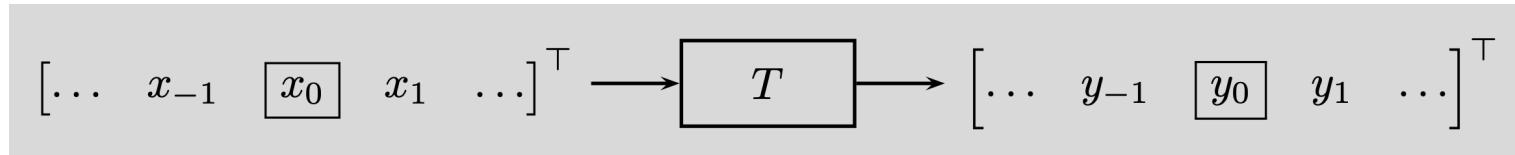
$$T(\alpha x + \beta y) = \alpha T(x) + \beta T(y). \quad (3.32)$$

DEFINITION 3.4 (SHIFT-INVARIANT SYSTEM) A discrete-time system T is called *shift-invariant* when, for any integer k and input x ,

$$y = T(x) \Rightarrow y' = T(x'), \quad \text{where } x'_n = x_{n-k} \text{ and } y'_n = y_{n-k}. \quad (3.37)$$

Filtering

Systems



A discrete-time system

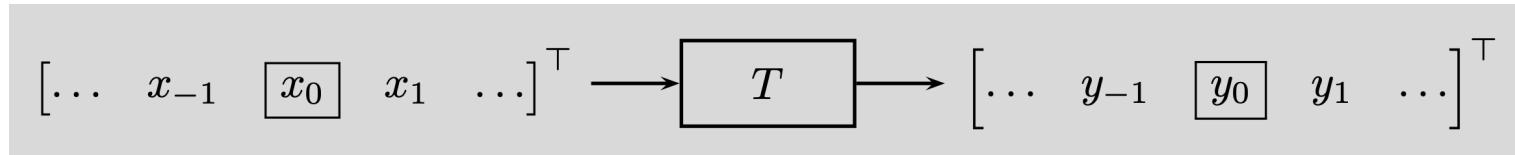
- Example

- The shift-by-1 operator, or delay $y_n = x_{n-1}$, $n \in \mathbb{Z}$

$$y = \begin{bmatrix} \vdots \\ y_{-1} \\ \boxed{y_0} \\ y_1 \\ \vdots \end{bmatrix} = \begin{bmatrix} \vdots \\ x_{-2} \\ \boxed{x_{-1}} \\ x_0 \\ \vdots \end{bmatrix} = \begin{bmatrix} \vdots & \vdots & \vdots & \vdots \\ \cdots & 0 & 0 & 0 & \cdots \\ \cdots & 1 & \boxed{0} & 0 & \cdots \\ \cdots & 0 & 1 & 0 & \cdots \\ \vdots & \vdots & \vdots & \vdots \end{bmatrix} \begin{bmatrix} \vdots \\ x_{-1} \\ \boxed{x_0} \\ x_1 \\ \vdots \end{bmatrix} \quad \text{LSI !}$$

Filtering

Systems



A discrete-time system

- Example

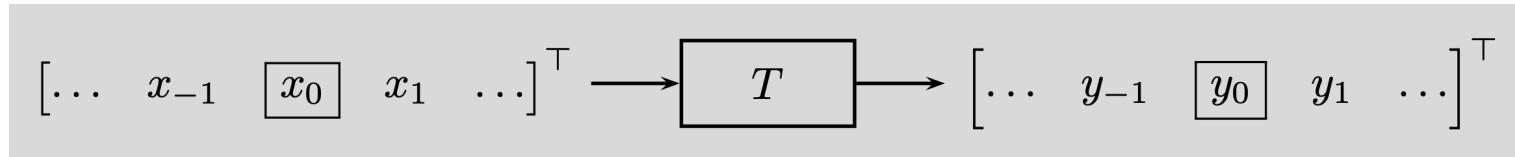
- The advance-by-1 operator $y_n = x_{n+1}, n \in \mathbb{Z}$

$$y = \begin{bmatrix} \vdots \\ y_{-1} \\ \boxed{y_0} \\ y_1 \\ \vdots \end{bmatrix} = \begin{bmatrix} \vdots \\ x_0 \\ \boxed{x_1} \\ x_2 \\ \vdots \end{bmatrix} = \begin{bmatrix} \vdots & \vdots & \vdots & \vdots \\ \cdots & 0 & 1 & 0 & \cdots \\ \cdots & 0 & \boxed{0} & 1 & \cdots \\ \cdots & 0 & 0 & 0 & \cdots \\ \vdots & \vdots & \vdots & \vdots \end{bmatrix} \begin{bmatrix} \vdots \\ x_{-1} \\ \boxed{x_0} \\ x_1 \\ \vdots \end{bmatrix}$$

LSI !

Filtering

Systems



A discrete-time system

- Example

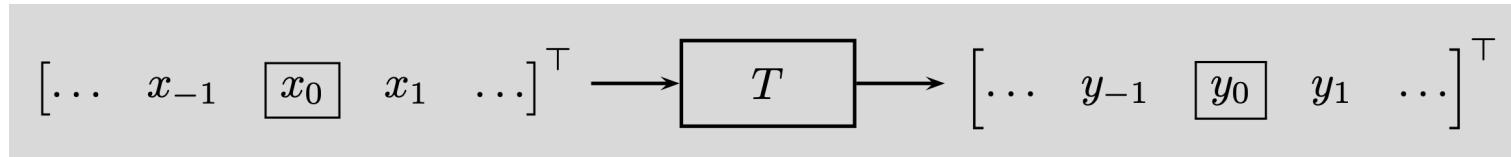
- Modulator $y_n = \alpha_n x_n, n \in \mathbb{Z}$

$$\begin{bmatrix} \vdots \\ y_{-1} \\ \boxed{y_0} \\ y_1 \\ \vdots \end{bmatrix} = \begin{bmatrix} \vdots \\ \alpha_{-1}x_{-1} \\ \boxed{\alpha_0x_0} \\ \alpha_1x_1 \\ \vdots \end{bmatrix} = \begin{bmatrix} \cdots & \alpha_{-1} & 0 & 0 & \cdots \\ \cdots & 0 & \boxed{\alpha_0} & 0 & \cdots \\ \cdots & 0 & 0 & \alpha_1 & \cdots \\ \vdots & \vdots & \vdots & \vdots & \vdots \end{bmatrix} \begin{bmatrix} \vdots \\ x_{-1} \\ \boxed{x_0} \\ x_1 \\ \vdots \end{bmatrix}$$

No LSI !

Filtering

Systems



A discrete-time system

- Example

- Averaging operators

$$y_n = \frac{1}{3}(x_{n-1} + x_n + x_{n+1}), \quad n \in \mathbb{Z}$$

$$\begin{bmatrix} \vdots \\ y_{-1} \\ \boxed{y_0} \\ y_1 \\ \vdots \end{bmatrix} = \frac{1}{3} \begin{bmatrix} \dots & \vdots & \vdots & \vdots & \vdots & \vdots \\ \dots & 1 & 1 & 1 & 0 & 0 & \dots \\ \dots & 0 & 1 & \boxed{1} & 1 & 0 & \dots \\ \dots & 0 & 0 & 1 & 1 & 1 & \dots \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \end{bmatrix} \begin{bmatrix} \vdots \\ x_{-2} \\ x_{-1} \\ \boxed{x_0} \\ x_1 \\ x_2 \\ \vdots \end{bmatrix}$$

LSI !

Filtering

Linear shift-invariant systems

DEFINITION 3.6 (IMPULSE RESPONSE) A sequence h is called the *impulse response* of LSI discrete-time system T when input δ produces output h .

$$x = \delta$$

$$[\dots \ x_{-1} \ \boxed{x_0} \ x_1 \ \dots]^{\top} \longrightarrow \boxed{T} \longrightarrow [\dots \ y_{-1} \ \boxed{y_0} \ y_1 \ \dots]^{\top}$$

Filtering

Linear shift-invariant systems

DEFINITION 3.7 (CONVOLUTION) The *convolution* between sequences h and x is defined as

$$(Hx)_n = (h * x)_n = \sum_{k \in \mathbb{Z}} x_k h_{n-k} = \sum_{k \in \mathbb{Z}} x_{n-k} h_k, \quad (3.61)$$

where H is called the *convolution operator* associated with h .

$$y = Tx = T \sum_{k \in \mathbb{Z}} x_k \delta_{n-k} \stackrel{(a)}{=} \sum_{k \in \mathbb{Z}} x_k T \delta_{n-k} \stackrel{(b)}{=} \sum_{k \in \mathbb{Z}} x_k h_{n-k} = h * x$$

↑ ↑ ↑
linearity shift invariance definition

Filtering

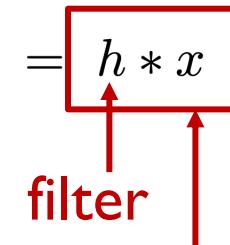
Linear shift-invariant systems

DEFINITION 3.7 (CONVOLUTION) The *convolution* between sequences h and x is defined as

$$(Hx)_n = (h * x)_n = \sum_{k \in \mathbb{Z}} x_k h_{n-k} = \sum_{k \in \mathbb{Z}} x_{n-k} h_k, \quad (3.61)$$

where H is called the *convolution operator* associated with h .

$$y = Tx = T \sum_{k \in \mathbb{Z}} x_k \delta_{n-k} \stackrel{(a)}{=} \sum_{k \in \mathbb{Z}} x_k T \delta_{n-k} \stackrel{(b)}{=} \sum_{k \in \mathbb{Z}} x_k h_{n-k} = \boxed{h * x}$$



The impulse response of a system is called a **filter**

Convolution with the impulse response is called **filtering**

Convolution is a mathematical way to implement a LSI system

Filtering

Linear shift-invariant systems

- Convolution

$$y = \begin{bmatrix} \vdots \\ y_{-2} \\ y_{-1} \\ \boxed{y_0} \\ y_1 \\ y_2 \\ \vdots \end{bmatrix} = \underbrace{\begin{bmatrix} \cdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \\ \cdots & h_0 & h_{-1} & h_{-2} & h_{-3} & h_{-4} & \cdots \\ \cdots & h_1 & h_0 & h_{-1} & h_{-2} & h_{-3} & \cdots \\ \cdots & h_2 & h_1 & \boxed{h_0} & h_{-1} & h_{-2} & \cdots \\ \cdots & h_3 & h_2 & h_1 & h_0 & h_{-1} & \cdots \\ \cdots & h_4 & h_3 & h_2 & h_1 & h_0 & \cdots \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \end{bmatrix}}_H \begin{bmatrix} \vdots \\ x_{-2} \\ x_{-1} \\ \boxed{x_0} \\ x_1 \\ x_2 \\ \vdots \end{bmatrix} = Hx$$

Filtering

Linear shift-invariant systems

- Properties of convolution

(i) *Connection to the inner product:*

$$(h * x)_n = \sum_{k \in \mathbb{Z}} x_k h_{n-k} = \langle x_k, h_{n-k}^* \rangle_k.$$

(ii) *Commutativity:*

$$h * x = x * h.$$

(iii) *Associativity:*

$$g * (h * x) = g * h * x = (g * h) * x.$$

(iv) *Deterministic autocorrelation:*

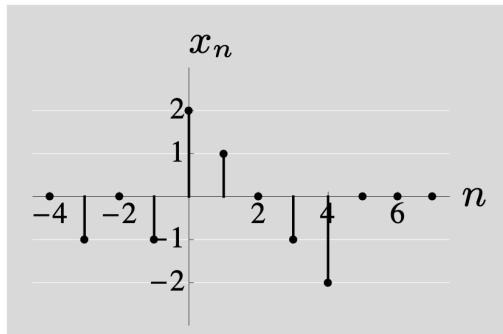
$$a_n = \sum_{k \in \mathbb{Z}} x_k x_{k-n}^* = x_n *_n x_{-n}^*.$$

(v) *Shifting:* For any $k \in \mathbb{Z}$,

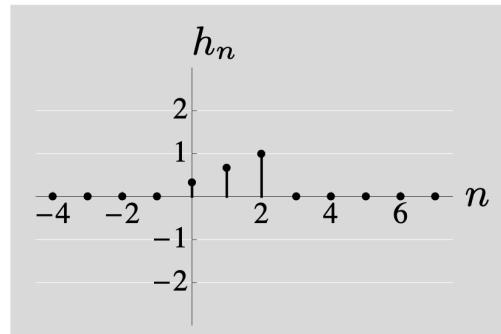
$$x_n *_n \delta_{n-k} = x_{n-k}.$$

Filtering

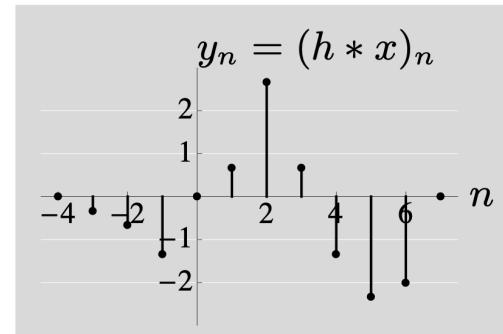
Linear shift-invariant systems



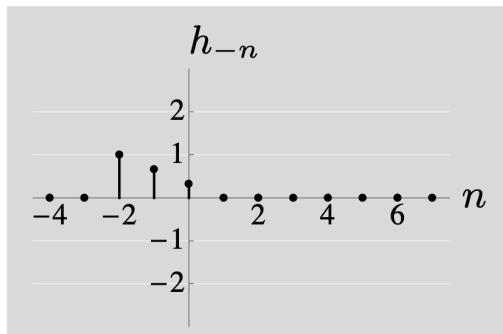
(a) Sequence.



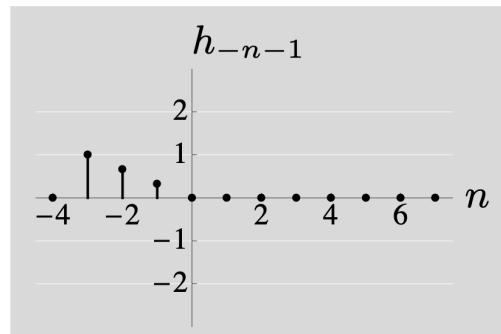
(b) Impulse response.



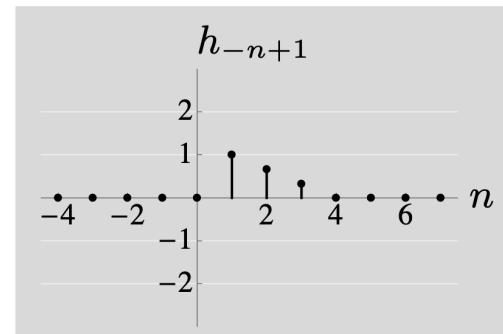
(c) Convolution.



(d)



(e)



(f)

Image filtering

Intro to filtering

More maths about filtering

Image filter design

Image filter design

DEFINITION 3.6 (IMPULSE RESPONSE) A sequence h is called the *impulse response* of LSI discrete-time system T when input δ produces output h .

$$x = \delta$$

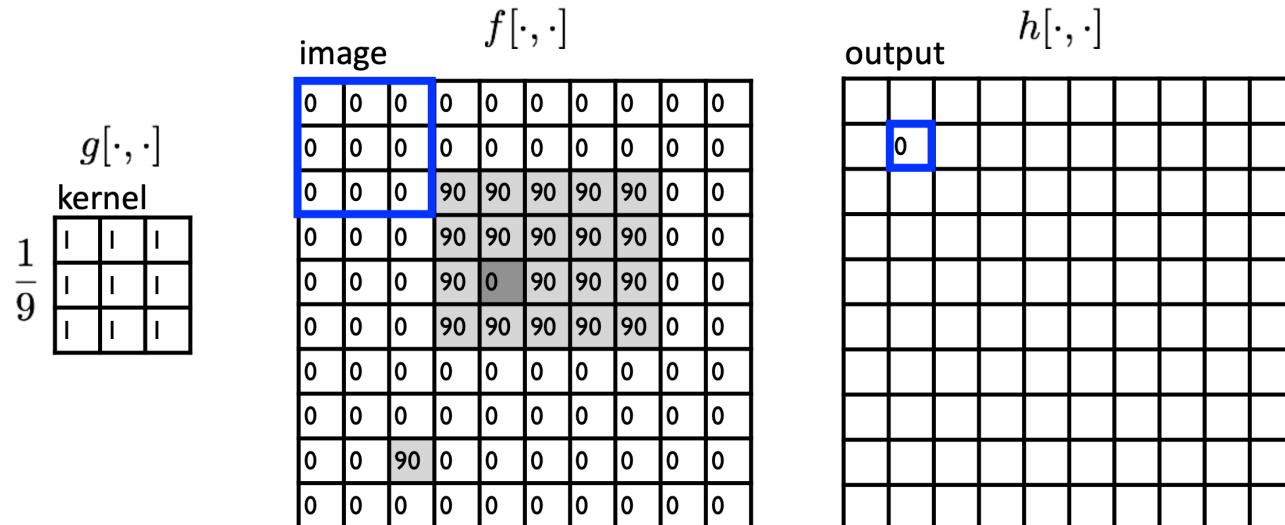
The diagram illustrates the convolution operation between an input image and a filter. On the left, a black square with a single white pixel in the center represents the input image x . In the middle, a small 3x3 grid of gray pixels represents the filter h , with the central pixel being white. To the right of the filter is an equals sign (=). To the right of the equals sign is a black square containing a single white 3x3 pixel block, representing the output image y .

Image filter design

DEFINITION 3.7 (CONVOLUTION) The *convolution* between sequences h and x is defined as

$$(Hx)_n = (h * x)_n = \sum_{k \in \mathbb{Z}} x_k h_{n-k} = \sum_{k \in \mathbb{Z}} x_{n-k} h_k, \quad (3.61)$$

where H is called the *convolution operator* associated with h .

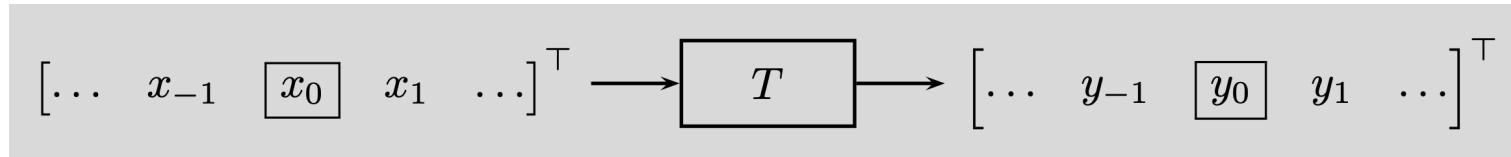


$$h[m, n] = \sum_{k, l} g[k, l] f[m + k, n + l]$$

output k, l filter image (signal)

Filtering

Systems



A discrete-time system

- Example

- Averaging operators

$$y_n = \frac{1}{3}(x_{n-1} + x_n + x_{n+1}), \quad n \in \mathbb{Z}$$

$$\begin{bmatrix} \vdots \\ y_{-1} \\ \boxed{y_0} \\ y_1 \\ \vdots \end{bmatrix} = \frac{1}{3} \begin{bmatrix} \dots & \vdots & \vdots & \vdots & \vdots & \vdots \\ \dots & 1 & 1 & 1 & 0 & 0 & \dots \\ \dots & 0 & 1 & \boxed{1} & 1 & 0 & \dots \\ \dots & 0 & 0 & 1 & 1 & 1 & \dots \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \end{bmatrix} \begin{bmatrix} \vdots \\ x_{-2} \\ x_{-1} \\ \boxed{x_0} \\ x_1 \\ x_2 \\ \vdots \end{bmatrix}$$

Image filter design

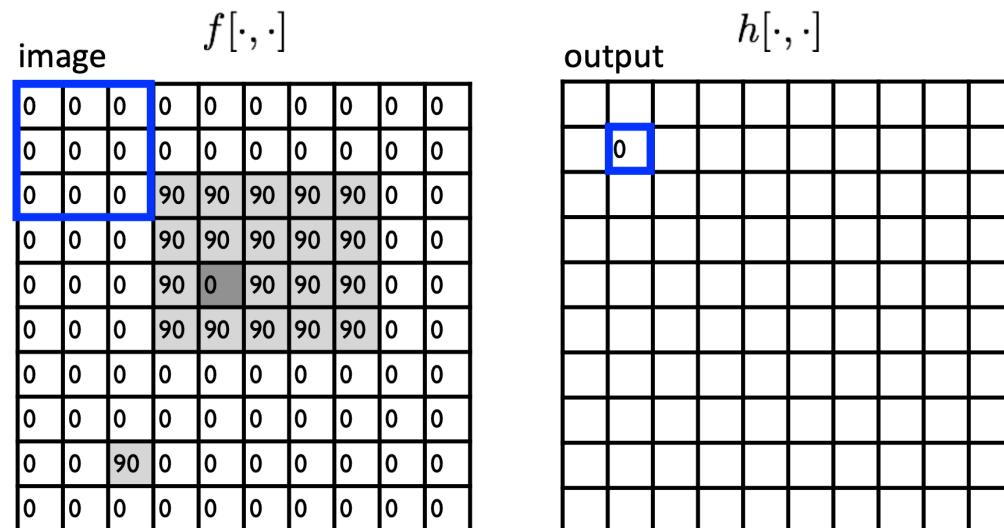
DEFINITION 3.7 (CONVOLUTION) The *convolution* between sequences h and x is defined as

$$(Hx)_n = (h * x)_n = \sum_{k \in \mathbb{Z}} x_k h_{n-k} = \sum_{k \in \mathbb{Z}} x_{n-k} h_k, \quad (3.61)$$

where H is called the *convolution operator* associated with h .

image filter:
averaging operator

$g[\cdot, \cdot]$
kernel
 $\frac{1}{9}$
 $\begin{array}{|c|c|c|} \hline | & | & | \\ \hline | & | & | \\ \hline | & | & | \\ \hline \end{array}$



$$h[m, n] = \sum_{k, l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

Image filter design

DEFINITION 3.7 (CONVOLUTION) The *convolution* between sequences h and x is defined as

$$(Hx)_n = (h * x)_n = \sum_{k \in \mathbb{Z}} x_k h_{n-k} = \sum_{k \in \mathbb{Z}} x_{n-k} h_k, \quad (3.61)$$

where H is called the *convolution operator* associated with h .

image filter:
averaging operator

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

Image filter design

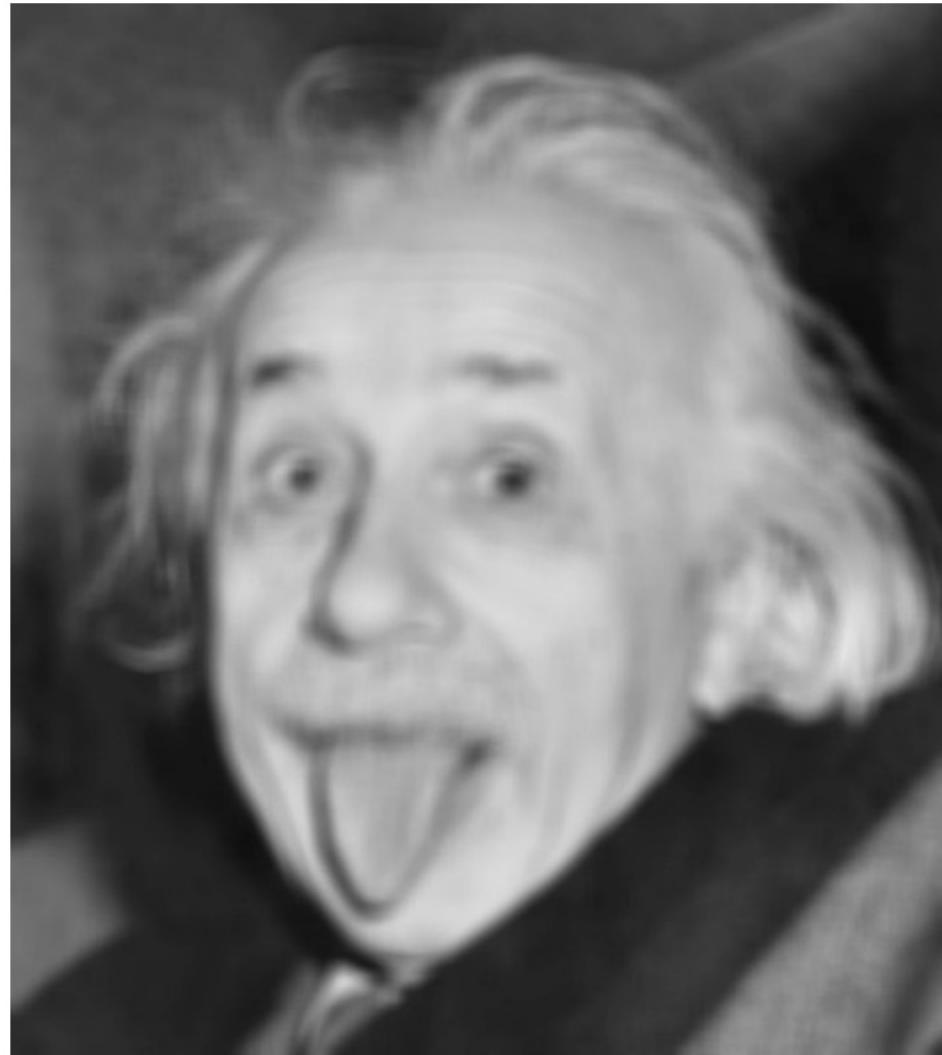
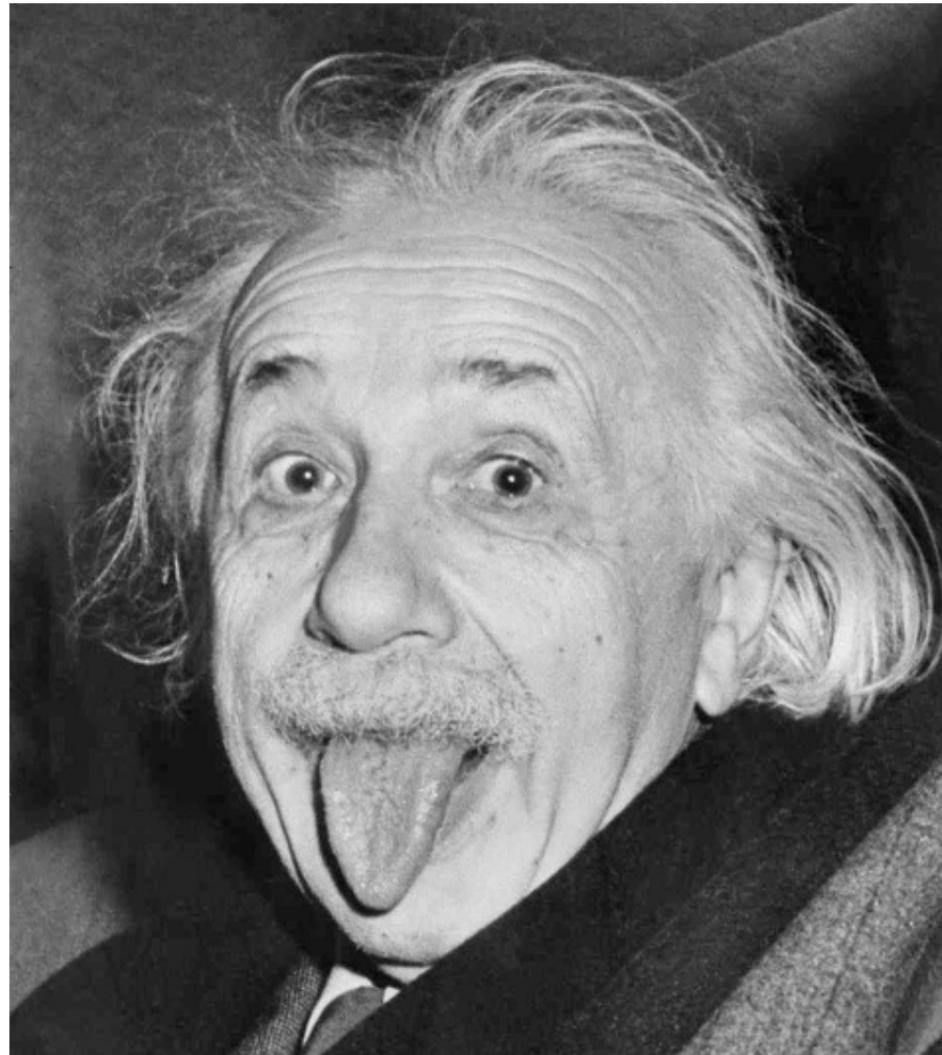


Image filter design



Image filter design

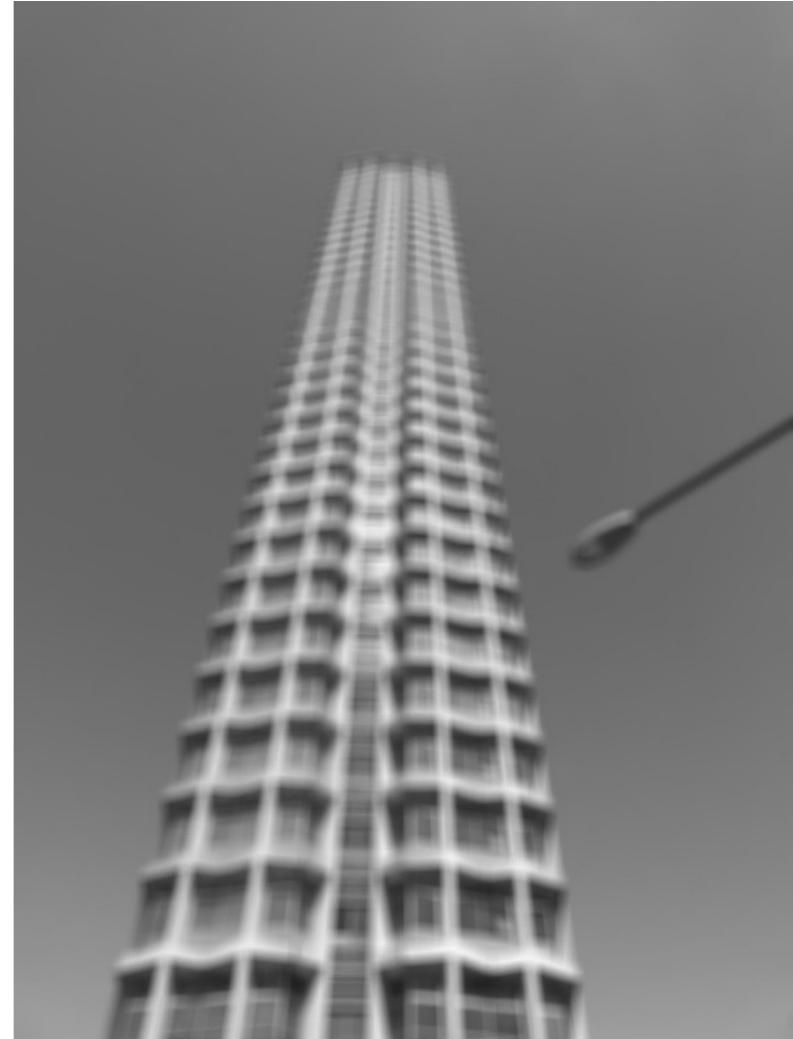


Image filter design

A 2D filter is separable if it can be written as the product of a “column” and a “row”.

$$\begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 1 \\ \hline 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline \end{array}$$

column

row

2D convolution with a separable filter is equivalent to two 1D convolutions (with the "column" and "row" filters).

If the image has $M \times M$ pixels and the filter kernel has size $N \times N$:

What is the cost of convolution with a non-separable filter?

$M^2 \times N^2$

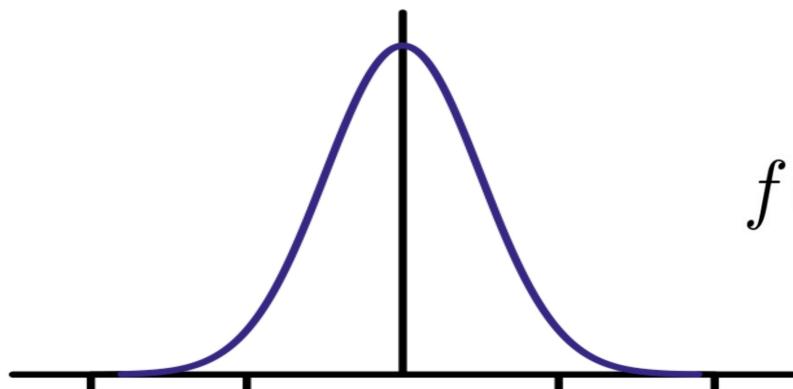
What is the cost of convolution with a separable filter?

$2 \times N \times M^2$

Image filter design

Gaussian filter

- kernel values sampled from the 2D Gaussian function



$$f(i, j) = \frac{1}{2\pi\sigma^2} e^{-\frac{i^2+j^2}{2\sigma^2}}$$

weight falls off with distance from center pixel

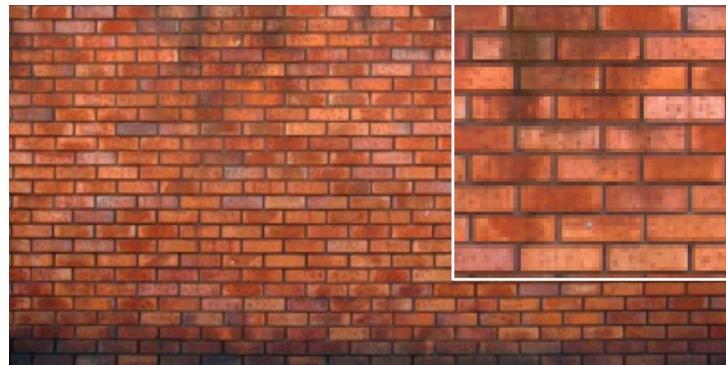
$\frac{1}{16}$

1	2	1
2	4	2
1	2	1

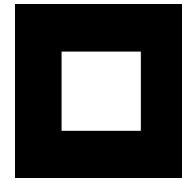
Is this a separable filter? YES !

Image filter design

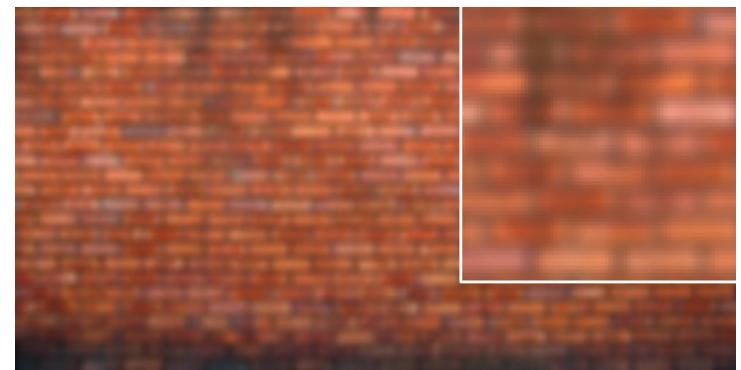
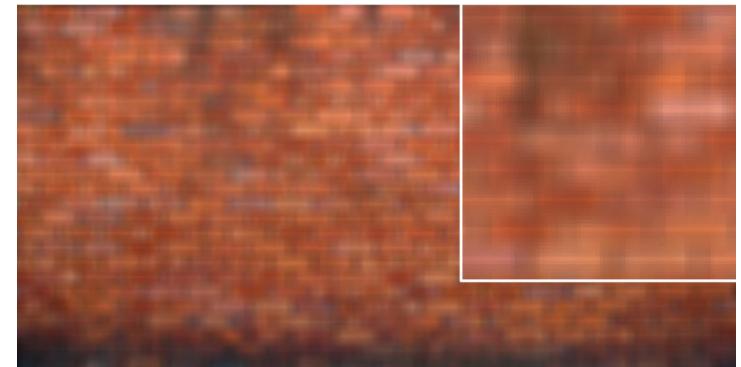
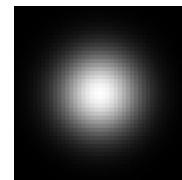
Gaussian filter



averaging operator



filtering



Gaussian filter

Smoothing!

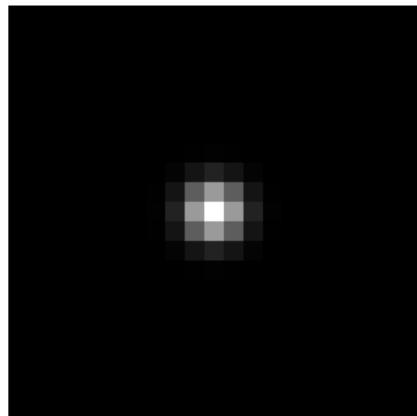
Image filter design

Gaussian filter

$$f(i, j) = \frac{1}{2\pi\sigma^2} e^{-\frac{i^2+j^2}{2\sigma^2}}$$

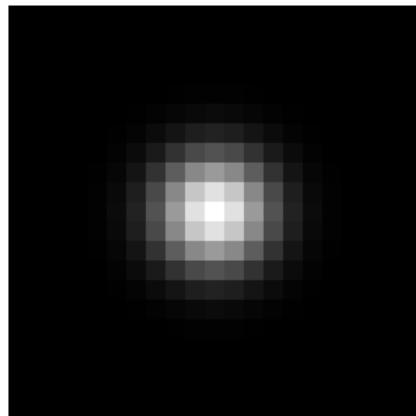
$\sigma = 1$

filter = 21x21



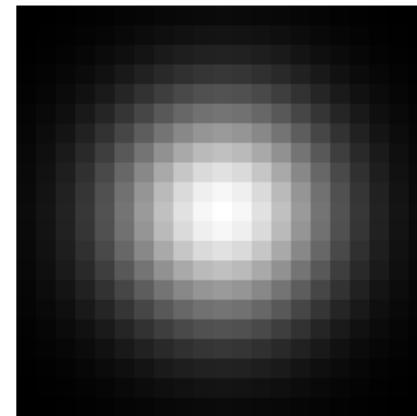
$\sigma = 2$

filter = 21x21



$\sigma = 4$

filter = 21x21



$\sigma = 8$

filter = 21x21

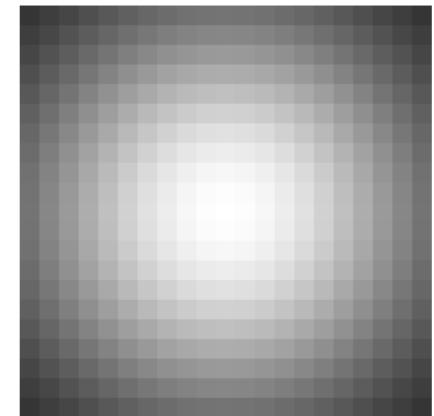


Image filter design

Gaussian filter

$$\sigma = 2$$

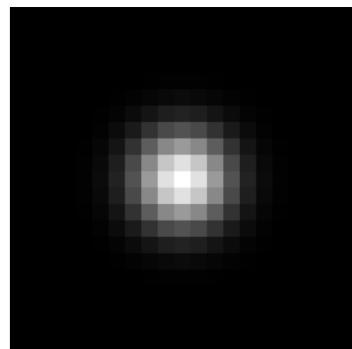


Image filter design

Gaussian filter

$$\sigma = 4$$

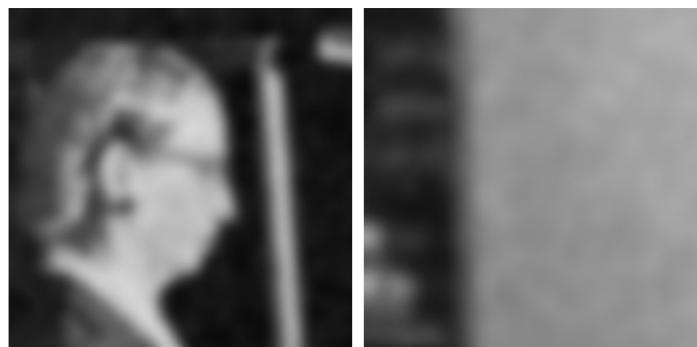
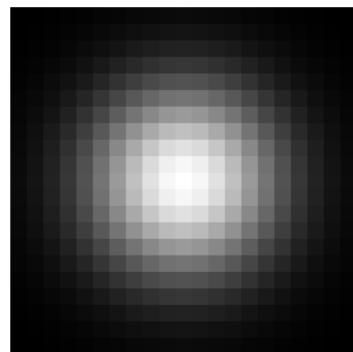


Image filter design

Gaussian filter

$$\sigma = 8$$

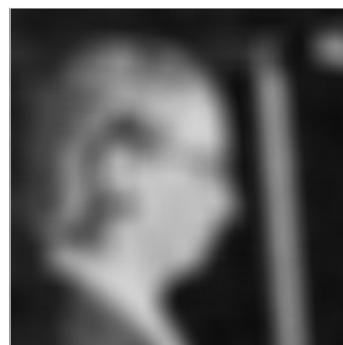
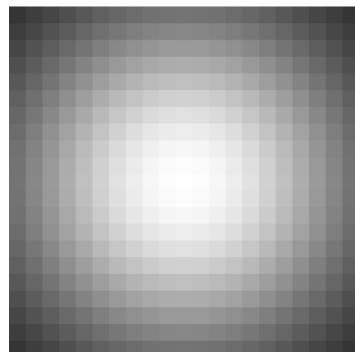


Image filter design

Gaussian filter

$$\sigma = 8$$

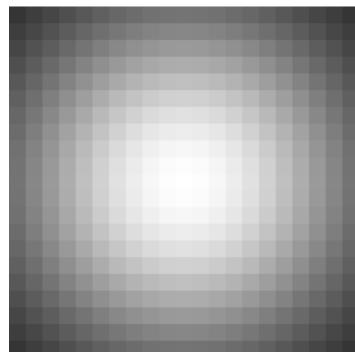


Image filter design

Sharpening (high-pass)

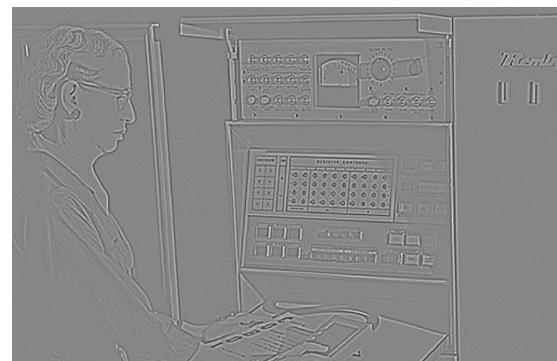
image



smoothed



details



=

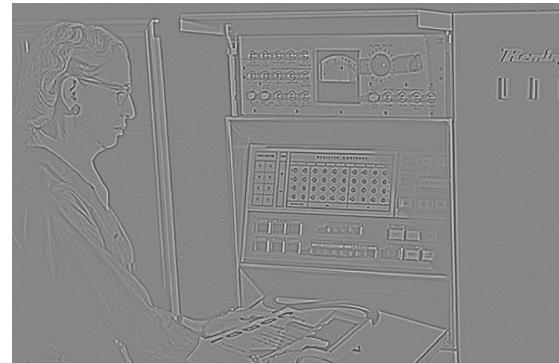
Image filter design

Sharpening (high-pass)

image



details



$+ \alpha$

sharpened

=



$\alpha = 1$

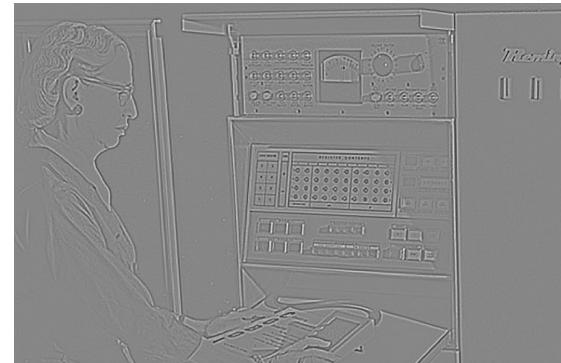
Image filter design

Sharpening (high-pass)

image



details



$+ \alpha$

sharpened

=



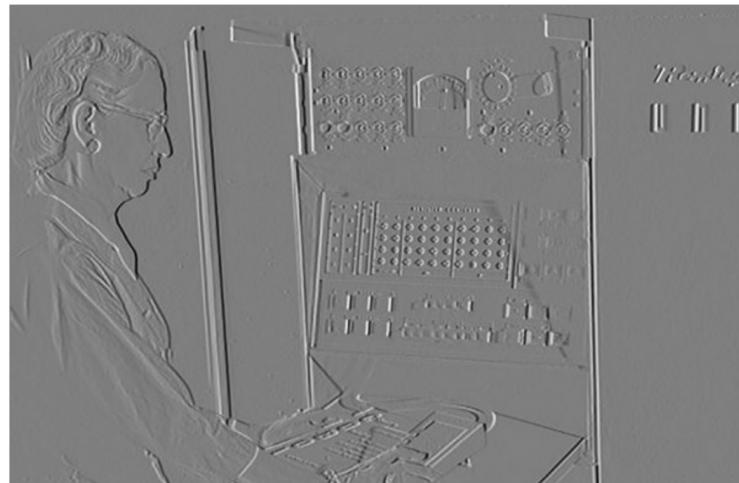
$\alpha = 10$

Image filter design

Sharpening (high-pass)

$$\begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline \end{array}$$

Dx



$$\begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline \end{array}$$

Dy

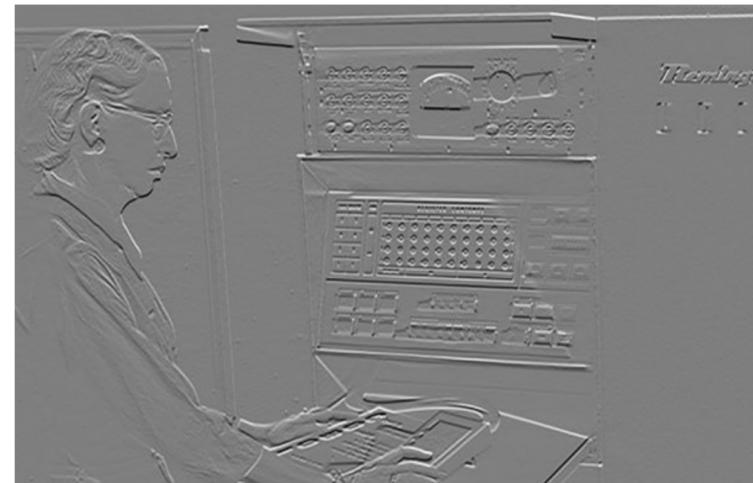


Image filter design

Gabor filter

$$g(x, y; \lambda, \theta, \psi, \sigma, \gamma) = \exp\left(-\frac{x'^2 + \gamma^2 y'^2}{2\sigma^2}\right) \exp\left(i\left(2\pi\frac{x'}{\lambda} + \psi\right)\right)$$

where $x' = x \cos \theta + y \sin \theta$ and $y' = -x \sin \theta + y \cos \theta$.

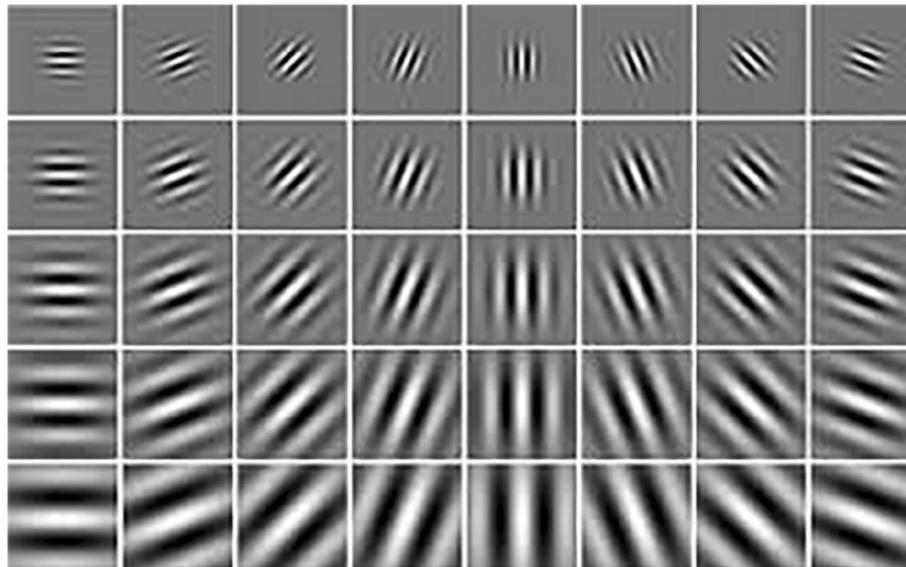


Image filter design

Convolutional neural network

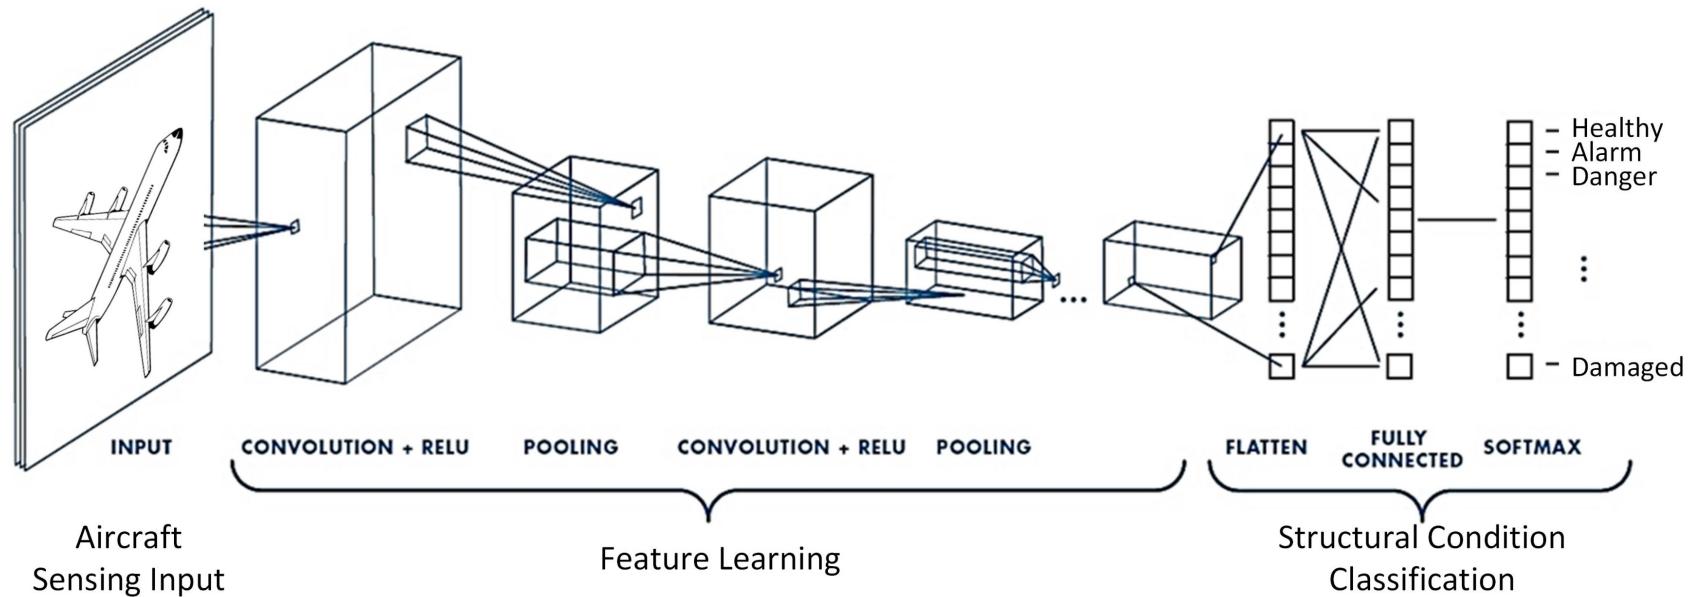
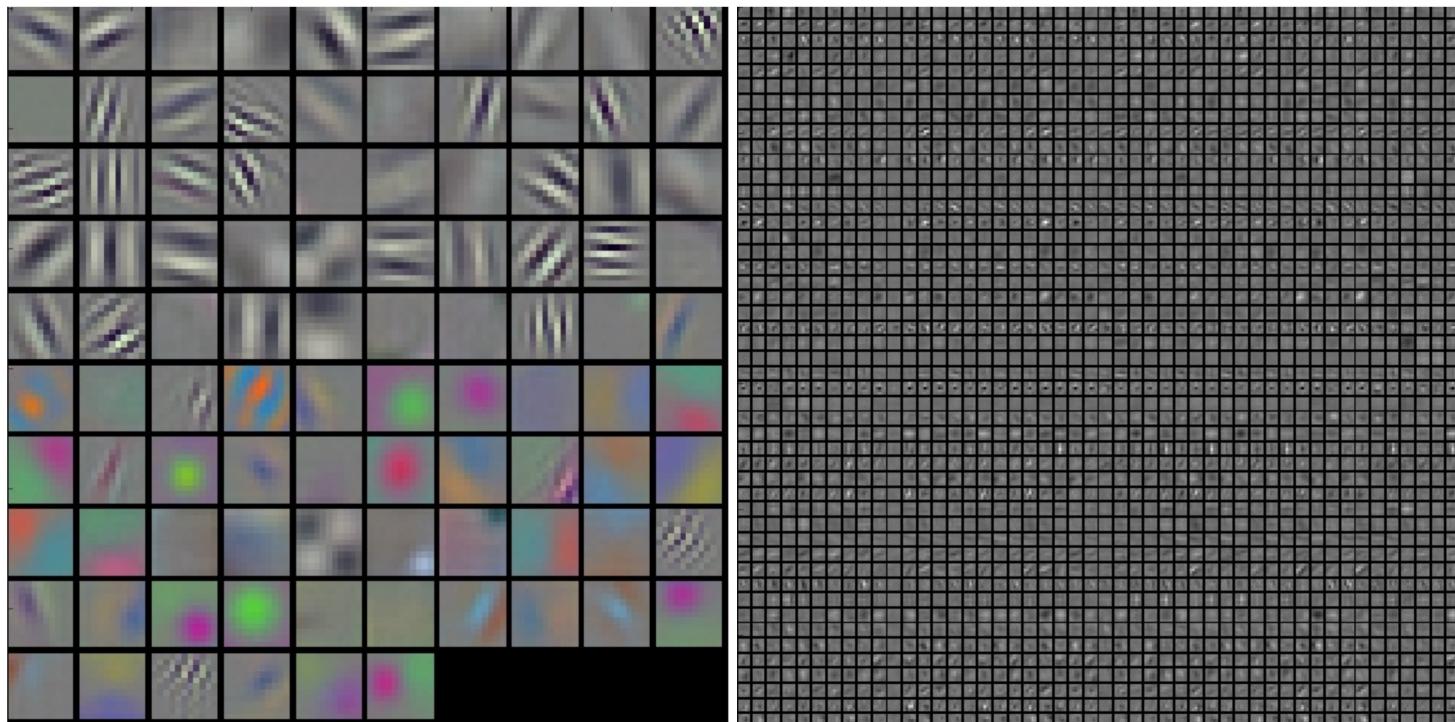


Image filter design

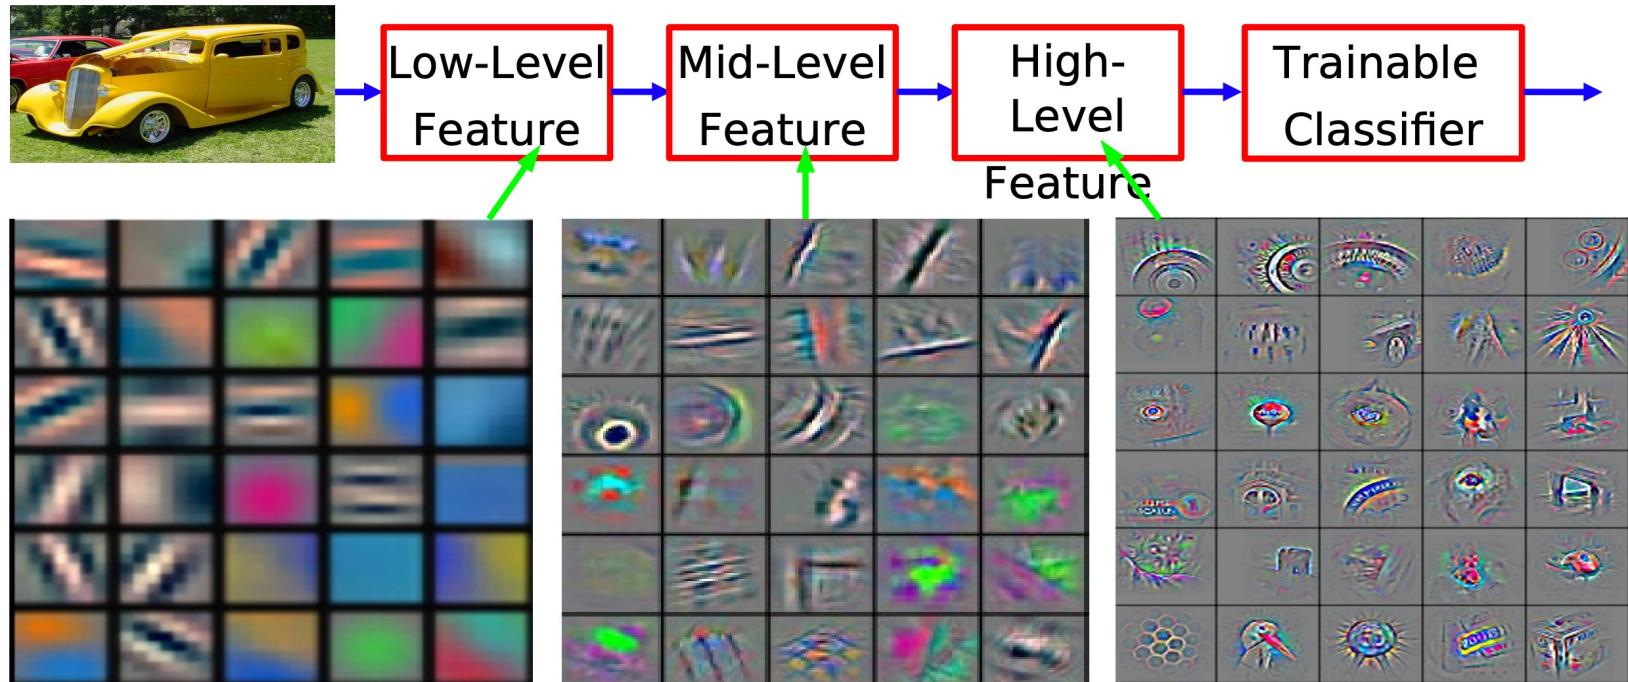
Convolutional neural network



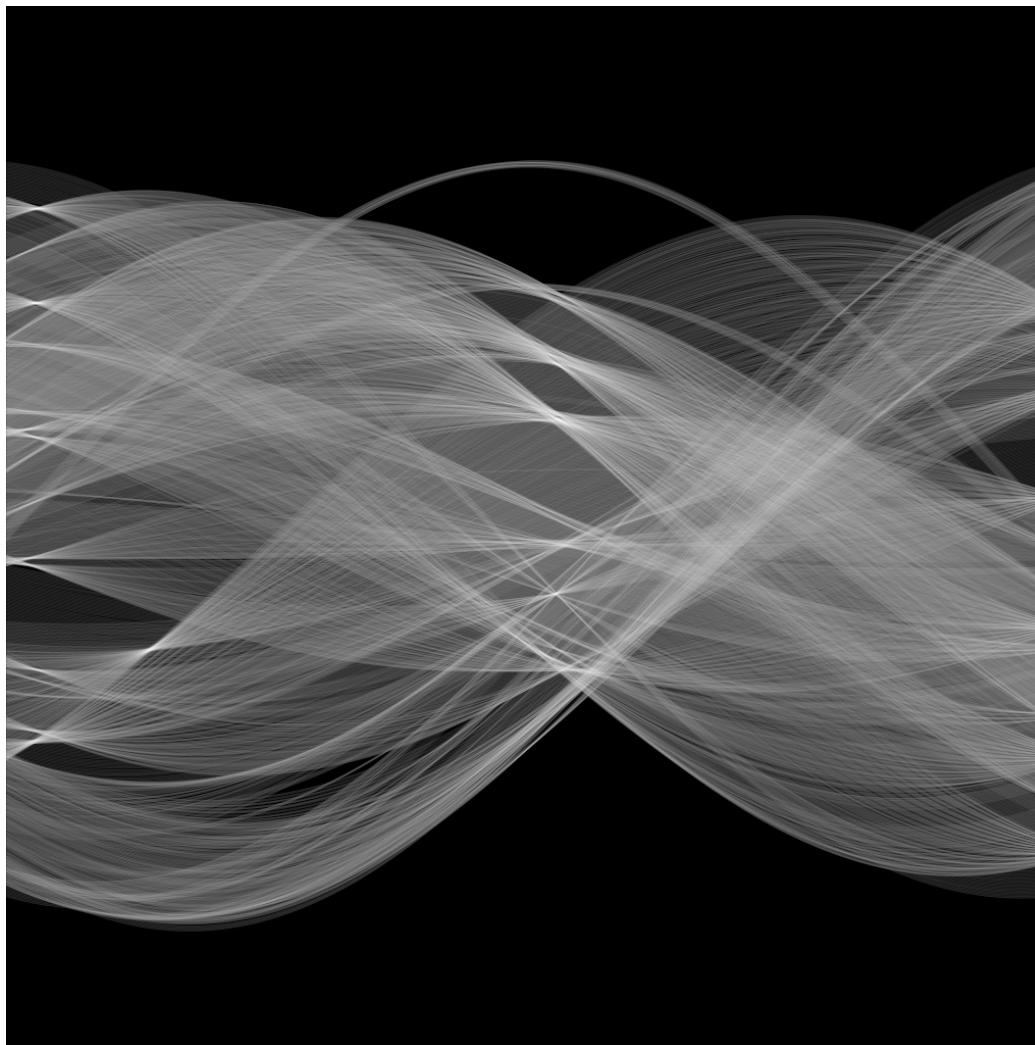
Typical-looking filters on the first CONV layer (left), and the 2nd CONV layer (right) of a trained AlexNet.

Image filter design

Convolutional neural network



Next lecture: Line detection



Thank you very much!

sihengc@sjtu.edu.cn